

[PLACEHOLDER TITLE]

Starting from the past: transient initialization and the input-history problem in soil carbon models

Lorenzo Menichetti [co-authors TBD]

Draft outline — June 17, 2026

Abstract

[PLACEHOLDER ABSTRACT.] One paragraph: steady-state initialization is a near-universal but rarely examined convention; an equilibrium start is already at its saturated end-state, so it cannot reproduce an accumulating sink that is only now beginning to saturate — the very dynamic observed in Finnish forest soils. We test a transient alternative across six models under one fair frame; it captures the accumulation-then-saturation as a proof of concept, but relocates the binding uncertainty onto litter inputs — especially their unknown history — pointing to historical, interdisciplinary reconstruction as the real frontier.

Contents

I	Opening	1
1	Introduction	1
1.1	The convention: steady state at t_0	1
1.1.1	What the assumption actually claims	1
1.1.1.1	Why it is attractive	1
1.1.1.2	Why it is rarely examined	1
1.1.2	The non-equilibrium reality	2
1.1.2.1	An equilibrium start is already at the end	2
1.1.3	Managed landscapes are rarely at equilibrium	2
1.2	Why this matters	2
1.2.1	Scientific stakes	2
1.2.2	Policy stakes: the motivating case	2
1.2.2.1	The observed dynamic	2
1.3	Scope and contribution (preview)	2
II	Challenge	2
2	The gap and the question	2
2.1	The saturation an equilibrium start cannot show	2
2.1.1	Why this generalizes beyond Finland	3
2.1.2	Artefact or structure?	3

2.2	The methodological gap	3
2.2.1	The fixed-value trap	3
2.2.2	Calibrate it, do not assume it	3
2.3	Objectives and hypotheses	3
III	Action	3
3	Materials and methods	3
3.1	Study system and data	3
3.1.1	Plots and SOC campaigns	3
3.1.2	Litter inputs	3
3.1.3	Climate forcing	4
3.2	The model ladder	4
3.2.1	Pool-complexity axis	4
3.2.2	Climate-integration axis	4
3.3	Bayesian calibration: the enabling machinery	4
3.3.1	Why a Bayesian frame is the natural fit	4
3.3.2	Uncertainty that propagates, not a likelihood patch	4
3.4	Transient initialization (the protagonist)	4
3.4.1	Pushing the equilibrium assumption back in time	4
3.4.1.1	The fast/slow relaxation argument	4
3.4.2	Calibrating the historical anchor	4
3.5	A fair frame for intercomparison	4
3.5.1	Shared data, likelihood, forcing, filtering	5
3.5.2	Exact integration in every model	5
3.5.3	Homogenized priors: method, not number	5
3.5.4	Equal degrees of freedom	5
3.6	Calibration and inference	5
3.6.1	Error model	5
3.6.2	Sampler and diagnostics	5
3.7	Predictive and residual analysis	5
4	Results	5
4.1	Capturing accumulation-then-saturation	5
4.1.1	Trajectory shape vs. level	5
4.1.2	Match to campaign means	5
4.2	Structural differences under the fair frame	5
4.2.1	Where simple beats complex	5
4.3	What the data can and cannot constrain	6
4.3.1	Convergence vs. identifiability	6
4.3.2	Input multiplier behaviour	6
4.4	Residual structure	6
IV	Resolution	6

5	Discussion	6
5.1	Capturing saturation requires remembering the accumulation	6
5.1.1	Proof of concept, not a verdict on Finland	6
5.1.2	Generalization to the model class	6
5.2	The binding uncertainty is the inputs (the σ_{input} reveal)	6
5.2.1	Historical inputs: the unrecoverable past	6
5.2.2	Belowground inputs and other ramifications	6
5.3	A short excursus: why this class of models is non-identifiable	7
5.3.1	The structural argument	7
5.3.1.1	Sketch	7
5.3.1.2	Signature in practice	7
5.3.1.3	Why it is not a defect to be “fixed”	7
5.4	The over-prediction is a property of the likelihood	7
5.5	Limitations	7
5.5.1	The coarse pre-observation phase	7
5.5.2	The flat-vs-rising litter tension	7
5.6	A managed landscape out of equilibrium: Finnish forest history	7
5.6.1	From exploitation to managed growth	7
5.6.2	The imprint on litter inputs and stocks	7
5.6.3	Why this is the rule, not a Finnish peculiarity	8
5.7	Outlook: toward a historical perspective on inputs	8
5.7.1	An interdisciplinary, source-diverse effort	8
5.7.2	Initialization as a calibratable, evidence-based step	8
6	Conclusions	8
6.1	The experiment succeeded	8
6.2	The Bayesian approach made it possible	8
6.3	Refinement needs historical data	8

Part I

Opening

1 Introduction

1.1 The convention: steady state at t_0

A near-universal default in SOC modelling. State it plainly as common practice, not yet as a problem.

1.1.1 What the assumption actually claims

That the system’s past equals its present equilibrium (inputs $\approx$ outputs at the start). Make the hidden content of the convention explicit.

1.1.1.1 Why it is attractive

Convenient, analytically clean, needs no historical forcing. The reasons it became default.

1.1.1.2 Why it is rarely examined

It is invisible precisely because it is standard. Set up the reader to recognize their own practice.

1.1.2 The non-equilibrium reality

Real ecosystems are transient; they carry the imprint of disturbance, land-use and productivity history.

1.1.2.1 An equilibrium start is already at the end

A steady-state start sits at zero net change — by construction already at its saturated end-state. It forecloses the accumulation phase before the simulation begins. Foreshadow the capture-failure without resolving it.

1.1.3 Managed landscapes are rarely at equilibrium

Land-use and management histories keep soils in transient states; equilibrium is the exception, not the rule. Flag Finnish forest history as the concrete case that justifies the violated assumption — stated generally here, examined in detail in the Discussion.

1.2 Why this matters

1.2.1 Scientific stakes

Initialization choices propagate into every downstream projection and into model intercomparison.

1.2.2 Policy stakes: the motivating case

Finland reports its forest soil as a sink; climate policy leans on it. Introduce Finland as the instance, not the subject.

1.2.2.1 The observed dynamic

Accumulation (VMI8→Biosoil), then recent hints of saturation (Komeetta, under review). A sink caught in the act of saturating. State the phenomenon the model must reproduce.

1.3 Scope and contribution (preview)

One short paragraph previewing the move: a transient initialization, tested fairly across a model ladder. Keep it a promise, not the argument.

Part II

Challenge

2 The gap and the question

2.1 The saturation an equilibrium start cannot show

The triggering insight: to reproduce a sink that is saturating, the model must first reproduce the accumulation it is the tail of — and an equilibrium start, already saturated, cannot. This is what

motivated transient init.

2.1.1 Why this generalizes beyond Finland

Any non-equilibrium system started at steady state forecloses its own transient. Lift the problem off the case study.

2.1.2 Artefact or structure?

Do divergent saturation projections reflect genuine model structure, or the initialization and calibration convention? The question the fair frame will answer.

2.2 The methodological gap

No established, calibratable way to initialize a transient run when historical forcing data are absent.

2.2.1 The fixed-value trap

Treating the unknown initial state as a single fixed number (equilibrium, or a guessed value) hides a strong, untested assumption inside the result.

2.2.2 Calibrate it, do not assume it

The alternative: make the historical start a parameter with a prior and let the data and its uncertainty speak. Plant the thread; pay it off in Methods.

2.3 Objectives and hypotheses

H1 Transient initialization lets the models reproduce the accumulation-then-saturation a steady-state start cannot. [placeholder]

H2 Residual model differences under a fair frame are structural and modest. [placeholder]

H3 The binding uncertainty migrates to inputs, not kinetics. [placeholder]

Part III

Action

3 Materials and methods

3.1 Study system and data

3.1.1 Plots and SOC campaigns

Finnish NFI plots; three temporally separated SOC campaigns as targets. Brief.

3.1.2 Litter inputs

Plot-level AWEN-fractionated inputs; note the population-mean rise then plateau (load-bearing later).

3.1.3 Climate forcing

Gridded daily climate; per-model temporal aggregation.

3.2 The model ladder

Six models on two crossed axes (pool complexity; climate integration), so no conclusion rests on one structure.

3.2.1 Pool-complexity axis

SP1 → TP2 → TP3 → Yasso07. One sentence each. [placeholder]

3.2.2 Climate-integration axis

Yasso07 → Yasso15 → Yasso20. One sentence each. [placeholder]

3.3 Bayesian calibration: the enabling machinery

The inference framework that lets the unknown start be calibrated rather than assumed. Keep it brief and matter-of-fact; its role is enabling, not headline.

3.3.1 Why a Bayesian frame is the natural fit

Priors encode what little is known about the past; the posterior carries uncertainty forward instead of collapsing it to a point.

3.3.2 Uncertainty that propagates, not a likelihood patch

σ_{init} is a prior on the historical state; its spread flows into both the initial condition and the forward dynamics. Distinguish from a mere error term added at the end.

3.4 Transient initialization (the protagonist)

3.4.1 Pushing the equilibrium assumption back in time

Initialize at analytical steady state in the past, run transiently forward to t_0 so the imprint dissipates and the model arrives at t_0 still accumulating.

3.4.1.1 The fast/slow relaxation argument

Fast pools fully relax; only the slow humus pool retains memory — it “remembers” the equilibrium start. Why the pre-run works, and where its memory lives.

3.4.2 Calibrating the historical anchor

σ_{init} as a Bayesian-calibrated strength on the historical start, not a fixed choice — the concrete place the thread meets the protagonist.

3.5 A fair frame for intercomparison

The design principle: any residual difference must be attributable to structure alone.

3.5.1 Shared data, likelihood, forcing, filtering

Brief enumeration.

3.5.2 Exact integration in every model

All six integrate their linear system exactly; differences are structural, not numerical.

3.5.3 Homogenized priors: method, not number

Tiered scheme; identical in construction/scale/DoF, not in raw values.

3.5.4 Equal degrees of freedom

All transfer fractions free in every Yasso variant. One sentence on why.

3.6 Calibration and inference

3.6.1 Error model

Multiplicative-normal likelihood. Flag the asymmetry briefly (paid back in Discussion).

3.6.2 Sampler and diagnostics

MCMC settings; convergence and identifiability diagnostics used.

3.7 Predictive and residual analysis

Posterior-predictive trajectories; residual analysis against site covariates.

4 Results

4.1 Capturing accumulation-then-saturation

The headline: transiently-initialized models reproduce the rise and its incipient bend-over; the equilibrium-started flat/declining shape was the start, not the structure.

4.1.1 Trajectory shape vs. level

Shape change, not just offset. [placeholder figure]

4.1.2 Match to campaign means

2006 and 2024 means tracked. [placeholder]

4.2 Structural differences under the fair frame

Real but modest skill spread across the six. [placeholder table]

4.2.1 Where simple beats complex

TP2/TP3 track shape best; stiffer Yasso family worst. Brief.

4.3 What the data can and cannot constrain

Set up the Discussion: kinetics weakly constrained; inputs do the work.

4.3.1 Convergence vs. identifiability

Good R^2 coexisting with non-identified fractions. State the apparent paradox, resolve in Discussion.

4.3.2 Input multiplier behaviour

σ_{input} does different work in different models. The hinge result.

4.4 Residual structure

RF importance; stand basal area as the leading, diffuse, shared driver. [placeholder figure]

Part IV

Resolution

5 Discussion

5.1 Capturing saturation requires remembering the accumulation

Resolve the Challenge: the difficulty was the equilibrium start, not the soil; a transient start lets every model represent the observed dynamic. The central payoff.

5.1.1 Proof of concept, not a verdict on Finland

What the result does and does not license. Calibrate the claim.

5.1.2 Generalization to the model class

Wherever steady-state init meets a non-equilibrium system, it forecloses the transient the data demand. Lift off Finland.

5.2 The binding uncertainty is the inputs (the σ_{input} reveal)

The data cannot resolve kinetics; the input multiplier absorbs what structure cannot. So the frontier is inputs, not soil chemistry.

5.2.1 Historical inputs: the unrecoverable past

No forcing data pre-observation; σ_{init} stands in for an entire history, held in the slow pool that remembers it. The deepest uncertainty.

5.2.2 Belowground inputs and other ramifications

Root/rhizodeposition inputs poorly known; a further axis of the same problem.

5.3 A short excursus: why this class of models is non-identifiable

5.3.1 The structural argument

First-order kinetics (and richer topologies) let parameters interact so the data constrain net fluxes out of a pool, not the individual splits.

5.3.1.1 Sketch

For a linear pool system the observable stocks depend on the parameters only through lumped combinations; distinct parameter sets give identical predictions (parameter-space ridges). Two equations of placeholder math. [placeholder]

5.3.1.2 Signature in practice

Prior-pinning, near-perfect anti-correlation ridges between same-pool fractions, low information gain (KL of posterior vs prior). The Bayesian frame is what lets us see this rather than hide it in a fitted point. Tie to Results.

5.3.1.3 Why it is not a defect to be “fixed”

It is intrinsic to the model class and the data content (few observations per plot); it bounds what calibration can ever claim.

5.4 The over-prediction is a property of the likelihood

Slight, systematic over-prediction follows from the multiplicative-normal error model, not from biogeochemistry. Only the bias is likelihood-sensitive; rankings are not.

5.5 Limitations

5.5.1 The coarse pre-observation phase

Linear litter interpolation under constant climate; high, deliberate uncertainty.

5.5.2 The flat-vs-rising litter tension

Plot-level inputs vs. national reconstructions; present data cannot settle it. State openly.

5.6 A managed landscape out of equilibrium: Finnish forest history

The concrete history behind the violated equilibrium assumption: why Finnish forest soils were accumulating, not at steady state, across the observation period. Pays off the Intro’s general statement.

5.6.1 From exploitation to managed growth

A sketch of the management eras — heavy historical use, then regeneration and intensified management, rising growing stock through the 20th century. [placeholder — sources to add]

5.6.2 The imprint on litter inputs and stocks

Rising biomass and shifting management make the input history non-stationary — precisely the signal an equilibrium start erases. Connect to the σ_{input} reveal and the flat-vs-rising litter tension.

5.6.3 Why this is the rule, not a Finnish peculiarity

Most managed and recovering landscapes carry comparable histories; the lesson generalizes. Tie back to the Intro’s general statement.

5.7 Outlook: toward a historical perspective on inputs

The forward call. The real next step is reconstructing historical (and belowground) input regimes.

5.7.1 An interdisciplinary, source-diverse effort

Reconstruction needs evidence beyond the usual datasets — grey literature and non-traditional historical sources. Frame as a research program.

5.7.2 Initialization as a calibratable, evidence-based step

Replace the equilibrium default with priors grounded in reconstructed history, calibrated under uncertainty. The thread closes here: better history → sharper prior → a start that is assumed less and inferred more. The generalizable recommendation.

6 Conclusions

6.1 The experiment succeeded

Transient, calibrated initialization let a model ladder reproduce the observed accumulation and its incipient saturation — the dynamic an equilibrium start forecloses. State the proof of concept plainly.

6.2 The Bayesian approach made it possible

Treating the historical start as a calibrated parameter, with uncertainty propagated into state and dynamics, is what turned the idea into a working implementation. Credit the machinery without overclaiming novelty.

6.3 Refinement needs historical data

The approach can go much further once the pre-observation history is reconstructed from data rather than left to calibration — the open frontier, and the natural successor study. Close here.

Notes for collapsing (remove before submission)

- **Parts** = OCAR acts; will dissolve into standard Intro/Methods/Results/Discussion at cleanup.
- Many heading levels are intentional scaffolding; merge aggressively once the story is fixed.
- Keep the protagonist (initialization problem) constant; do not let σ_{input} or non-identifiability take over the spine.