

HIKET — the storyline, v2

skeleton for discussion: headers, one-line arguments, figures

Version of September 2, 2026. Numbers keyed to RUN_IDs 20260831_1624* (arm B). Deliberately not prose — the point is to argue about the ARC.

How to read this. Every block is a claim we intend to make, its evidence, and nothing else. ▶ **name** marks a figure that exists; ★ **TO BUILD: name** marks one that does not yet. Green boxes are *tests we could run*; amber boxes are *things that could sink the claim*.

Contents

1	The arc in one paragraph	2
I	Act I — Opening	2
2	A sink caught saturating	2
3	The convention that cannot see it	2
II	Act II — Challenge	3
4	To show the saturation, the model must remember it was rising	3
III	Act III — Action	3
5	Transient initialisation under a fair frame	3
IV	Act IV — Resolution, part 1: it works	3
6	The models reach the observations	3
V	Act V — Thread A, restated	4
7	Fit does not discriminate structure; dynamics do	4
7.1	In-sample, structure is invisible	4
7.2	Out of sample, structure is nearly everything	4
7.3	The applied consequence	4
VI	Act VI — Thread B, the new centre	4

8	The calibration wants more input <i>and</i> faster kinetics	5
8.1	Why the surprise is the method's own doing	5
8.2	Inputs	5
8.3	Kinetics	5
8.4	Why these are ONE claim, not two	6
8.5	The independent witness (not on the ridge)	6
VII	Act VII — Consequences, and how to test them	6
9	If the ensemble sits elsewhere on the ridge, what changes?	6
9.1	The hypothesis	6
9.2	The proper test: published Yasso vs. ours, forward only	7
9.2.1	The published arm agrees with the data. That is the whole point.	7
9.2.2	Pre-registration	8
9.2.3	What we actually publish: one figure in the main text	8
9.3	Stronger versions, if this one pays off	10
VIII	Act VIII — Thread C	10
10	What the data can and cannot constrain	10
IX	Act IX — Landing	10
11	The policy deliverable	10
X	Act X — Outlook	10
12	Where the recoverable signal hides	11
XI	Appendix	11
13	Figure inventory	11

1 The arc in one paragraph

Finnish forest soils are accumulating carbon and the accumulation is beginning to level off. A model started at steady state cannot show that, because a steady state is already the end of it. Started instead from a calibrated past, all six models reach the observed levels and two of them reproduce the rate exactly — but doing so requires *more litter input and faster turnover* than the published parameterisations carry. That displacement is the finding, and it is not free: it should change how the soil responds to a warming climate. **The new work is to test whether it does.**

Part I

Act I — Opening

2 A sink caught saturating

2.0.0.1 The observed dynamic. On the balanced plot set the mean stock runs $64.5 \rightarrow 71.5 \rightarrow 73.6 \text{ tC ha}^{-1}$ across the three campaigns: +7.0 then +2.1. Accumulation, bending over. ► [F3_mean_soc](#)

2.0.0.2 The rate is what an inventory reports. $+0.259 \pm 0.040 \text{ tC ha}^{-1} \text{ yr}^{-1}$ over the full span. ► [F5_stock_change](#)

2.0.0.3 Why it is politically live. The Finnish GHG inventory turns on whether this sink persists or saturates. ► [F10b_forest_history](#)

3 The convention that cannot see it

3.0.0.1 Steady-state initialisation is near-universal and rarely examined. Lehtonen 2016, Palosuo 2008, Peltoniemi 2004 — our own group’s lineage, to be framed as evolving, not attacked.

3.0.0.2 An equilibrium start is already at the end. It cannot produce a rise, so it cannot produce a rise that is levelling off.

Part II

Act II — Challenge

4 To show the saturation, the model must remember it was rising

4.0.0.1 The soil carries a memory the model is told to forget. Slow pools integrate a century of inputs; setting them to equilibrium erases exactly the transient the data show.

4.0.0.2 Managed landscapes are rarely at equilibrium. Finland is the vivid case, not a special one: exploitation, then a century of managed regrowth. ► [F10b_forest_history](#)

Part III

Act III — Action

5 Transient initialisation under a fair frame

5.0.0.1 The unknown start becomes a calibrated parameter. Each draw spins up from 1917; σ_{init} is the ratio of 1917 to 1985 litter input. Uncertainty propagates rather than being patched. ► [F4_initialization](#)

5.0.0.2 Fair means equal *information*, not equal parameter count. Yasso does not fit its decomposition rates, so neither do the simple models: rates externally anchored (ICBM), fractions free. Otherwise structure and calibration asymmetry are confounded.

5.0.0.3 One likelihood, and it does not pretend the plots are independent. Log-normal with a fixed correlated error structure (region, plot, campaign), total $\sigma = 0.800$ partitioned, not added. ★ TO BUILD: [schematic of the variance partition — optional](#)

Part IV

Act IV — Resolution, part 1: it works

6 The models reach the observations

6.0.0.1 Levels. No model sits more than 3.5 tC ha^{-1} from any campaign mean. The old “ 26 tC ha^{-1} too high at 1985” caveat is gone. ► [F3_mean_soc](#)

6.0.0.2 The rate. Against observed $+0.259$: Yasso07 $+0.258$, Yasso15 $+0.250$. TP2/TP3 about a third strong; SP1/Yasso20 about half weak. The ensemble brackets the answer. ► [F5_stock_change](#)

6.0.0.3 The three-way split — new, and better than “all six work”. Yasso07/Yasso15 rise then flatten (*saturation reproduced*); TP2/TP3 rise and keep rising (*no saturation*); SP1/Yasso20 overshoot then decline. Structure is legible in the shape. ★ TO BUILD: [F3b — trajectory shape, three groups, explicitly grouped](#)

The 2006–2024 interval discriminates far less than it looks: observed $+0.117 \pm 0.066$, CI includes zero. Do not lean on “two models get the sign wrong” over eighteen years.

Part V

Act V — Thread A, restated

7 Fit does not discriminate structure; dynamics do

This replaces “does complexity improve the answer?”. The old answer (“no, and the most complex fits worst”) is no longer true — Yasso15 has the best calibration R^2 and the best holdout RMSE. The defensible and more useful claim is the one below.

7.1 In-sample, structure is invisible

7.1.0.1 Skill does not climb the pool ladder. Calibration R^2 0.016–0.021; holdout $R^2 \leq 0.0012$ in all six; holdout RMSE spans 3.6%. One pool and five pools are indistinguishable. ► [T1_model_summary](#) ► [F6b_skill_vs_complexity](#)

7.1.0.2 The ordering does not follow structure at all. Best and worst holdout RMSE belong to a five-pool and a three-pool model.

7.2 Out of sample, structure is nearly everything

7.2.0.1 The forecast diverges where the calibration agreed. ► [F4_initialization](#) (panel b)

7.2.0.2 Climate sensitivity differs $\sim 2\times$ across the family. Equilibrium SOC change at $+2^\circ\text{C}$: Yasso07 -21.7% , Yasso20 -13.7% , Yasso15 -9.6% . Between-model range is $3.1\times$ the mean within-model band. ► [climate_sensitivity_sweep](#)

Measured on an earlier run and as a *demonstration* (step change, single-exponential on bulk MRT). Must be recomputed on the current calibration before it goes in a figure.

7.3 The applied consequence

7.3.0.1 Model choice is a forecast decision, not a fit decision. Selecting an inventory model on calibration skill selects on a quantity that carries almost no information about what the model will project. This is the practical message of the paper.

7.3.0.2 So the deliverable is an ensemble, not a winner. An estimate plus its structural uncertainty.

Part VI

Act VI — Thread B, the new centre

8 The calibration wants more input *and* faster kinetics

This is the turn, and it should land as a surprise. The paper sets out to fix an initialisation convention. The expected ending is the tidy one: the convention was the problem, we replaced it, the models work. That ending is *half* true — and the other half is the interesting part. Making the models reproduce the observed trajectory requires a soil that is **fed more and turns over faster** than any of the published parameterisations describe. A methods fix walks into a claim about the models themselves.

8.1 Why the surprise is the method's own doing

8.1.0.1 A steady-state start hides exactly this. At equilibrium $C = J/k$: only the *product* matters, so any input-and-turnover pair with the right product reproduces the stock. Nothing in the fit can object to the published pair.

8.1.0.2 A transient start adds timing, and timing separates them. Once the model must *travel* from 1917 through three observed campaigns, the rate of approach matters, and that is set by the turnover time on its own. The constraint the convention discarded is the one that moves the ensemble off the published point.

8.1.0.3 So the twist is not incidental to the method — it is caused by it. The initialisation fix is what makes the parameter displacement visible. That is why the two halves belong in one paper.

Do not overstate the separation. The timing constraint *weakens* the ridge, it does not break it: MRT and σ_{input} still correlate -0.64 to -0.73 in the Yasso family, and the displacement remains a joint one. The claim is that the ridge acquired a slope, not that the two became independently identified.

8.2 Inputs

8.2.0.1 The multiplier has a physical referent. $\sigma_{\text{input}} \times \bar{J}$ is an effective litter flux: $3.19\text{--}4.34 \text{ tC ha}^{-1} \text{ yr}^{-1}$ across the six, i.e. 18–61% above the national total including understorey (~ 2.7).

► [SX_input_vs_literature](#)

8.2.0.2 Independent corroboration exists. Liski 2006 — same country, period, model family and inventory — implies a comparable correction. ► [F1_inputs](#)

8.3 Kinetics

8.3.0.1 Turnover is displaced from published values. Intrinsic MRT $24.6 / 27.1 / 22.6 \text{ yr}$ (Yasso07/15/20) against published points $33.5 / 30.4 / 19.0$. ► [F12_mrt_yasso](#)

8.3.0.2 Not one phenomenon. Yasso07’s whole gap is its *single* climate modifier (MRT = $\text{MRT}_{\text{ref}}/\xi$ exactly); Yasso15/20 carry three pool-specific modifiers and cannot use that lever. Do not discuss the family’s gaps as one.

8.4 Why these are ONE claim, not two

8.4.0.1 The level pins only the product. $\text{level} \sim \text{MRT} \times \sigma_{\text{input}} \times \bar{J}$. Faster-with-more and slower-with-less fit identically; the posterior rides that ridge (Yasso07/15/20 corr $-0.73/-0.68/-0.64$; the simple models stronger still). ► [F14_mrt_ridge](#) ► [S13_mrt_ridge_benchmark](#)

The discipline this imposes. “Inputs too low” and “turnover too slow” are two coordinates of *one* displacement. Report it once. Claiming both as independent findings would double-count, and CLAUDE.md flags this in three places.

8.4.0.2 So what *is* the claim? That the ensemble’s best-fitting position on the ridge lies at higher input and faster turnover than the published parameterisations sit. That is a single, well-posed statement about location — and it is testable, because the two ends of the ridge are *not* equivalent in everything.

8.5 The independent witness (not on the ridge)

8.5.0.1 The models leave an input-shaped signal in their residuals. They explain $\sim 2\%$ of SOC variance; a random forest explains 44–65% of what they leave. ► [F11_rf_heatmap](#)

8.5.0.2 And the covariate is the same in all six. Stand basal area — a productivity proxy, not a kinetic one. This does not ride the ridge, so it is genuinely separate evidence that the limitation is on the input side.

Part VII

Act VII — Consequences, and how to test them

9 If the ensemble sits elsewhere on the ridge, what changes?

The scientific core of the next phase. A displacement that changed nothing observable would be bookkeeping. The bet is that it changes the *climate response* — which would make the position on the ridge an empirically decidable question rather than a prior-driven one, and would turn the surprise of Act VI from a curiosity into a consequence for anyone projecting these soils.

9.1 The hypothesis

9.1.0.1 Equal fit, unequal reactivity. Two calibrations indistinguishable in fit, sitting at different points on the $\text{MRT} \times \sigma_{\text{input}}$ ridge, project *different* sensitivity to warming.

9.1.0.2 Why it should be so. Temperature acts on the *rate constants*. A faster-turnover solution has more of its carbon in pools that feel ξ on a decadal timescale, so the same warming moves more carbon. The input side responds through productivity instead — a different pathway, and one held fixed in these projections.

9.1.0.3 Already partly visible. Fit explains only 10–18% of forecast spread: 82–90% of the projected divergence is among draws the data cannot rank.

9.2 The proper test: published Yasso vs. ours, forward only

Rejected first: arm A vs arm B. The σ_{input} anchor arms look like a two-calibration experiment but are not a useful one. They share the same model freedom and the same data, and their optima barely separate — MRT 22.1 vs 24.6 yr, σ_{input} within 5–8%. That is a nudge along the ridge, not the two ends of it.

9.2.0.1 The real contrast is already published. The old Yasso parameterisation sits at *lower input and slower turnover*; our calibration sits at *higher input and faster turnover*. That is a genuine separation — Yasso07 MRT 33.5 with tree-litter-only input, against 24.6 with $\sigma_{\text{input}} = 1.59$: 36% in turnover and 59% in input, three to ten times the arms’.

9.2.0.2 Forward mode only — no re-calibration. Both parameter sets are pushed through the same engine, same plots, same litter series, same climate, same integrator. Nothing is fitted. This sidesteps the objection that a re-calibration lets every other parameter readjust, and it is cheap.

Design — settled, and the published arm is built. Two axes, one readout.

Axis 1 — parameters	published kinetics, σ_{input} refitted to the observed window ours (arm B posterior)
Axis 2 — climate	control (recycled observed, as now) warming (+2 °C, +5 °C)
Readout	<i>the response</i> : warming – control, per parameter set

σ_{init} is pinned to our posterior in **BOTH** arms, so the two differ in exactly two coordinates — turnover and input — and no initialisation difference is smuggled into a comparison that is about inputs versus kinetics.

9.2.1 The published arm agrees with the data. That is the whole point.

9.2.1.1 Holding published kinetics fixed, only the input multiplier is refitted. One parameter, scored on the three campaign means (balanced set). ▶ [doublechecks/published_kinetics_input_refit.R](#)

	published arm		our arm		contrast		misfit (pub/ours) rms log
	MRT	σ_{input}	MRT	σ_{input}	MRT	input	
Yasso07	33.47	1.249	24.57	1.585	$\times 0.73$	$\times 1.27$	0.020 / 0.029
Yasso15	29.97	1.358	27.11	1.538	$\times 0.90$	$\times 1.13$	0.023 / 0.034
Yasso20	22.10	1.708	22.56	1.659	$\times 1.02$	$\times 0.97$	0.029 / 0.035

9.2.1.2 No compromise was needed. The published arm matches the observed window *better than our calibration does* in all three models. So any divergence downstream is extrapolation, not level mismatch — which is exactly the condition the experiment required.

9.2.1.3 Both arms are on the same ridge. The product $\text{MRT} \times \sigma_{\text{input}}$ is conserved to 7%, 2.5% and 0.8% respectively. We are moving *along* each model’s ridge, not off it — which is what makes the contrast interpretable.

9.2.1.4 The contrast is graded, and Yasso20 is a negative control. Yasso07 is displaced hard ($\times 0.73$ in turnover, $\times 1.27$ in input), Yasso15 moderately, and **Yasso20 hardly at all** — its published parameterisation already sits where our calibration puts it. That is a gift for the design: if climate response tracks position on the ridge, Yasso07 must separate, Yasso15 less, and **Yasso20 must not separate at all**. A within-family null is a far stronger control than three replicates of one contrast.

Getting Yasso20 in required a fix, and the fix is documented. Published Yasso20 is calibrated with the AWEN transfer fractions *plus* p_H exhausting the pool budget **exactly** (median 1.000000; Yasso15 sits at 0.99939). Our Fortran `model_step` declares a matrix ill-conditioned when a column’s off-diagonal sum exceeds $0.9999 \times$ the diagonal — which is precisely this budget — and falls back to an explicit Euler step that diverges. The exact path is in fact stable *and fast* (1 ms) at 0.9998, and SOC converges smoothly as the budget rises (0.990 $\rightarrow$ 51.3; 0.998 $\rightarrow$ 52.5; 0.9998 $\rightarrow$ 52.75), so we clamp the budget at 0.9998, costing $\sim 0.02\%$. **The proper fix is in the Fortran** (raise the threshold; fall back only above 1, where the system genuinely creates mass), but that touches a `.so` shared by Yasso15 and Yasso20 and needs its own verification.

9.2.2 Pre-registration

9.2.2.1 Direction. Ours should be *more* reactive: faster turnover puts more carbon in pools that feel ξ on a decadal horizon, so the same warming moves more of it.

9.2.2.2 Where the effect should be largest. Yasso07, whose single ξ ties turnover and climate sensitivity to the same parameter (corr -0.58); smaller in Yasso15/20, where three pool-specific modifiers decouple them (-0.10 , -0.16).

9.2.2.3 What would falsify it. Responses that coincide within the posterior bands. That would say the ridge is degenerate in climate response as well as in fit — the position on it is then not decidable from projection, which is itself worth reporting.

$\sigma_{\text{input}} = 1$ on the published arm means *our* litter product taken as given, not the litter estimate Yasso was originally calibrated against. State it; it is the cleanest available choice but it is a choice.

9.2.3 What we actually publish: one figure in the main text

This is an important *side* story, not the spine. One figure in the manuscript, everything else in the appendix — and the main-text one leads with the **consequence**, not the mechanism.

9.2.3.1 F15, four panels (BUILT). Titles *describe* what is plotted; the reading belongs in the caption and the text, not in the panel heading.

9.2.3.2 (a) Modelled SOC, climate held stationary. All three Yassos, both arms (solid = ours, tight dash = published), 1985 $\rightarrow$ 2084, campaign means with CIs, rule at 2024.

9.2.3.3 (b) Modelled SOC, warming to SSP2-4.5. Same models, **same axes**, projection warmed on a linear ramp. (a) rises; (b) turns every model over into decline.

9.2.3.4 (c) Warming response at 2084, by arm. Boxplots over 200 draws.

9.2.3.5 (d) Difference between the arms. Boxplots with a zero line: only Yasso07 excludes zero.

DECIDED 2026-09-02 (Lorenzo). (i) The input anchor is Lehtonen & Heikkinen — the σ_{input} prior centre 1.08, i.e. what the run log calls arm B. Liski's 1.27 (arm A) is retained as a sensitivity, out of the draft, in the notes. (ii) **This side arm concludes about the difference in CLIMATE RESPONSE**, not about the predicted stock. Both are computed — the prediction gap has the same central values but every interval crosses zero, because absolute stock carries the full within-arm posterior spread — and the response is the claim because it is the paired, within-draw contrast that the data can actually speak to. Report both; lead with the response.

The result, and the words for it (Lorenzo, 2026-09-02). Yasso15 and Yasso20 prove **resilient to the input-anchor choice** — *not* “right”. Their pool-specific climate modifiers leave the carbon-holding pools (H and N, $\sim 60\%$ of the stock) responding to warming almost identically in both arms, while the arms differ mainly in the AWE modifier, and AWE holds $\sim 13\%$. Yasso07 has *one* modifier for every pool, so displacing it displaces the sensitivity of all the carbon at once. This is not a frontal attack on Yasso; it is a measured statement about which structure is robust to a calibration assumption the data cannot settle.

9.2.3.6 Measured mechanism (steady-state decomposition predicts the forward runs). $d \log \xi_H / +5^\circ\text{C}$ is $+0.338$ vs $+0.340$ for Yasso20's two arms — a 0.6% difference — and $+0.182$ vs $+0.165$ for Yasso15. Yasso07's single ξ goes $+0.308 \rightarrow +0.412$, a 34% change applied to every pool. Predicted response ratios $1.35/1.046/0.963$ against measured $1.41/1.031/0.934$.

9.2.3.7 Corollary worth stating in the paper. Bulk MRT is the *wrong* summary statistic for climate sensitivity in a multi-pool model with pool-specific modifiers: two parameterisations can differ 10% in MRT and be identical in warming response. The right statistic is the stock-weighted $\sum_i w_i d \log \xi_i / dT$.

9.2.3.8 Units and honesty. All panels in $tC ha^{-1}$: the absolute response carries both the kinetics and the input difference, which is the consequence a reader cares about. The relative version — σ_{input} -invariant, therefore pure kinetics — goes to the appendix. Yasso07’s published arm has no `.dat`, so it is a *line* where the others have bands, and must be drawn to look different rather than silently thinner.

9.2.3.9 Appendix A1 (BUILT). (a) response against each draw’s intrinsic MRT; (b) the same contrast in relative units; (c) all three SSP scenarios. Plus `T_A1`, how the published arm was constructed.

The collapse test came out as predicted, and it is now a measured result. A single common curve (response $\sim \log$ MRT) pooled over all models explains only $R^2 = 0.449$. Within models: Yasso07 **0.703** — its single ξ ties turnover and climate sensitivity to the same parameter — against Yasso15 0.209 and Yasso20 **0.068**, where three pool-specific modifiers decouple them. Visually the three clouds do not lie on one curve: Yasso20 at MRT ≈ 22 responds -8.8% while Yasso07 draws at the *same* MRT respond $\approx -10\%$. **Bulk MRT does not determine the warming response**, which is the evidence for the corollary below.

Why the collapse figure matters even in an appendix. It is what converts “two parameterisations differ” into “position on the ridge governs reactivity”. If the clouds do not collapse onto one curve, the effect is model identity rather than ridge position — a different claim, and one this design would not support.

★ TO BUILD: F15 (main text, 3 panels) + appendix: MRT collapse, all-model trajectories, arm table

9.3 Stronger versions, if this one pays off

9.3.0.1 A real scenario rather than a temperature offset. Replace the $+2/+5^\circ C$ construct with a proper climate projection, so the answer is a projection and not a sensitivity.

9.3.0.2 Let productivity respond too. Warming acts on inputs as well as on decomposition; here the input side is held fixed, which biases the comparison toward loss in *both* arms. Coupling it is the honest version and a much larger job.

9.3.0.3 Break the degeneracy from outside. Radiocarbon constrains turnover independently of input magnitude. The degeneracy is a property of the observation design, not of nature. Worth stating even if we cannot do it here.

The projections currently contain no climate change — the last twenty observed years are recycled. Everything above is therefore about *sensitivity*, computed separately, and the forecast figures are disequilibrium relaxation under stationary climate. The paper must say so.

Part VIII

Act VIII — Thread C

10 What the data can and cannot constrain

10.0.0.1 Convergence is not identifiability. $R\text{-hat} \leq 1.008$ everywhere, yet the transfer fractions are prior-pinned, ridge-correlated and carry < 1 nat. A clean sampler on a flat ridge. ► [F8b_fraction_ridges](#) ► [T2_fraction_ridges](#)

10.0.0.2 The data pin a level, and little else. Two observations per plot; no cross-plot skill. ► [F6_obs_vs_pred_holdout](#)

10.0.0.3 Report it, do not fix it. Non-identifiability here is a property of the observation design and belongs in the results.

Part IX

Act IX — Landing

11 The policy deliverable

11.0.0.1 An estimate plus its irreducible structural uncertainty. Not a single number, and not a selected winner.

11.0.0.2 Aggregate, not plot-level. These models capture the national trajectory; they do not rank plots. Say it early — otherwise “captures the dynamic” and “ $R^2 \approx 0$ ” read as a contradiction.

11.0.0.3 Maps as the landing figure. ► [NextGenC ensemble maps](#)

Part X

Act X — Outlook

12 Where the recoverable signal hides

12.0.0.1 Historical inputs. The deepest uncertainty; σ_{init} stands in for a century.

12.0.0.2 Understorey and belowground. The named gap in the litter product.

12.0.0.3 Stratified inputs. The fraction rises strongly south to north; one national multiplier cannot carry it.

12.0.0.4 Radiocarbon. The clean external constraint on turnover.

Part XI

Appendix

13 Figure inventory

13.0.0.1 Exist, need rebuilding on arm B. F1, F3, F4, F6, F6b, F11, F12, F14, S13, T1, T2 — all PNGs currently predate the current calibration.

13.0.0.2 Exist and are current. F5 (rebuilt 2026-09-02).

13.0.0.3 To build. F3b (shape groups); **F15** (main text, three panels: agreement $\rightarrow$ divergence $\rightarrow$ size of the disagreement, with Yasso20 as the visible null); appendix set (MRT collapse, all-model trajectories, arm-construction table, relative-response version).